

```

import os

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
from PIL import Image

import torch
from torch.utils.data import DataLoader, Dataset, Subset
from torchvision import transforms, datasets
from sklearn.model_selection import train_test_split
from tqdm import tqdm
import time
from sklearn.metrics import f1_score, accuracy_score

from google.colab import drive
drive.mount('/content/drive')
print(os.listdir('/content/drive/My Drive/Colab Notebooks/Aerial_Landscapes'))

Mounted at /content/drive
['Airport', 'Agriculture', 'City', 'Beach', 'Desert', 'Forest', 'Highway', 'Grassland', 'Lake', 'Mountain', 'Port', 'Parking', 'Residential', 'R

img_path = '/content/drive/My Drive/Colab Notebooks/Aerial_Landscapes'

transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(30),
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225])
])

dataset_path = img_path

dataset = datasets.ImageFolder(root=dataset_path, transform=transform)

print("类别名称:", dataset.classes)
print("数据集大小:", len(dataset))

data_loader = DataLoader(dataset, batch_size=32, shuffle=True, num_workers=4)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(torch.cuda.get_device_name())
torch.cuda.empty_cache()

类别名称: ['Agriculture', 'Airport', 'Beach', 'City', 'Desert', 'Forest', 'Grassland', 'Highway', 'Lake', 'Mountain', 'Parking', 'Port', 'Railway']
数据集大小: 12000
Tesla T4
/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes in total.
warnings.warn(

indices = np.arange(len(dataset))

y = np.array(dataset.targets)

train_idx, test_idx = train_test_split(indices, test_size=0.2, random_state=42, stratify=y)
print("train data:", len(train_idx))
print("test data:", len(test_idx))

train_dataset = Subset(dataset, train_idx)
test_dataset = Subset(dataset, test_idx)

batch_size = 32
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=16, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, num_workers=16, pin_memory=True)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

```

num_classes = len(dataset.classes)
model = timm.create_model('efficientnet_b0', pretrained=True, num_classes=num_classes)
model = model.to(device)

```

```

criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

```

```

num_epochs = 10
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    for images, labels in tqdm(train_loader, desc=f"Epoch {epoch+1}"):
        images = images.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item() * images.size(0)

    epoch_loss = running_loss / len(train_dataset)
    print(f"Epoch {epoch + 1}/{num_epochs} - Loss: {epoch_loss:.4f}")

    model.eval()
    correct, total = 0, 0
    y_true, y_pred = [], []
    with torch.no_grad():
        for images, labels in test_loader:
            images = images.to(device)
            labels = labels.to(device)
            outputs = model(images)
            preds = outputs.argmax(dim=1)
            correct += (preds == labels).sum().item()
            total += labels.size(0)
            y_true.extend(labels.cpu().tolist())
            y_pred.extend(preds.cpu().tolist())

    accuracy = correct / total
    print(f"Test Accuracy: {accuracy:.4f}")
    print("=== Classification Report ===")
    print(classification_report(
        y_true,
        y_pred,
        target_names=dataset.classes,
        digits=4
    ))

```



训练集样本数: 9600

测试集样本数: 2400

/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 16 worker processes in to warnings.warn(

Epoch 1: 100% ██████████ | 300/300 [00:56<00:00, 5.27it/s]Epoch 1/10 - Loss: 0.4906

Test Accuracy: 0.9408

=== Classification Report ===

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Agriculture | 0.9671 | 0.9187 | 0.9423 | 160 |
| Airport | 0.9346 | 0.8938 | 0.9137 | 160 |
| Beach | 0.9244 | 0.9938 | 0.9578 | 160 |
| City | 0.9455 | 0.9750 | 0.9600 | 160 |
| Desert | 0.9608 | 0.9187 | 0.9393 | 160 |
| Forest | 0.9675 | 0.9313 | 0.9490 | 160 |
| Grassland | 0.9235 | 0.9812 | 0.9515 | 160 |
| Highway | 0.9850 | 0.8187 | 0.8942 | 160 |
| Lake | 0.9603 | 0.9062 | 0.9325 | 160 |
| Mountain | 0.9728 | 0.8938 | 0.9316 | 160 |
| Parking | 0.9349 | 0.9875 | 0.9605 | 160 |
| Port | 0.9811 | 0.9750 | 0.9781 | 160 |
| Railway | 0.8112 | 0.9938 | 0.8933 | 160 |
| Residential | 0.9815 | 0.9938 | 0.9876 | 160 |
| River | 0.9085 | 0.9313 | 0.9198 | 160 |
| accuracy | | | 0.9408 | 2400 |
| macro avg | 0.9439 | 0.9408 | 0.9407 | 2400 |
| weighted avg | 0.9439 | 0.9408 | 0.9407 | 2400 |

Epoch 2: 0% | 0/300 [00:00<?, ?it/s]/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This l warnings.warn(

Epoch 2: 100% ██████████ | 300/300 [00:57<00:00, 5.18it/s]Epoch 2/10 - Loss: 0.1741

Test Accuracy: 0.9463

=== Classification Report ===


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Agriculture | 0.9444 | 0.9563 | 0.9503 | 160 |
| Airport | 0.9494 | 0.9375 | 0.9434 | 160 |
| Beach | 0.9693 | 0.9875 | 0.9783 | 160 |
| City | 0.8541 | 0.9875 | 0.9159 | 160 |
| Desert | 0.9865 | 0.9125 | 0.9481 | 160 |
| Forest | 0.9796 | 0.9000 | 0.9381 | 160 |
| Grassland | 0.8870 | 0.9812 | 0.9318 | 160 |
| Highway | 0.9812 | 0.9812 | 0.9812 | 160 |
| Lake | 0.9655 | 0.8750 | 0.9180 | 160 |
| Mountain | 0.9317 | 0.9375 | 0.9346 | 160 |
| Parking | 0.9576 | 0.9875 | 0.9723 | 160 |
| Port | 0.9737 | 0.9250 | 0.9487 | 160 |
| Railway | 0.9796 | 0.9000 | 0.9381 | 160 |
| Residential | 0.9815 | 0.9938 | 0.9876 | 160 |
| River | 0.8869 | 0.9313 | 0.9085 | 160 |
| accuracy | | | 0.9463 | 2400 |
| macro avg | 0.9485 | 0.9463 | 0.9463 | 2400 |
| weighted avg | 0.9485 | 0.9463 | 0.9463 | 2400 |

```

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
all_preds, all_labels = [], []
model.eval()
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        logits = model(images)
        preds = logits.argmax(dim=1)
        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

cm = confusion_matrix(all_labels, all_preds)
disp = ConfusionMatrixDisplay(cm, display_labels=dataset.classes)
disp.plot(cmap=plt.cm.Blues)
plt.title("Test Set Confusion Matrix")
plt.show()

```



开始借助 AI 编写或生成代码。

