

```

import os, torch, numpy as np, matplotlib.pyplot as plt, seaborn as sns
from torchvision import datasets, transforms, models
from torchvision.models import resnet50, ResNet50_Weights
from torch.utils.data import DataLoader, random_split
import torch.nn as nn
import torch.optim as optim
from torch.optim.lr_scheduler import StepLR
from sklearn.metrics import classification_report, confusion_matrix
import kagglehub

# Download dataset from Kaggle Hub
root_path = kagglehub.dataset_download("ankit1743/skyview-an-aerial-landscape-dataset")
data_path = os.path.join(root_path, "Aerial_Landscapes")

# Set device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Load pretrained ResNet-50 and modify final layer
weights = ResNet50_Weights.DEFAULT
model = resnet50(weights=weights)
model.fc = nn.Sequential(
    nn.Linear(model.fc.in_features, 512),
    nn.ReLU(),
    nn.Dropout(0.4),
    nn.Linear(512, 15)
)
model = model.to(device)

# Data augmentation and normalization
train_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(15),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])
test_transform = weights.transforms()

# Load dataset
full_dataset = datasets.ImageFolder(data_path, transform=train_transform)
train_size = int(0.8 * len(full_dataset))
test_size = len(full_dataset) - train_size
train_dataset, test_dataset = random_split(full_dataset, [train_size, test_size])
test_dataset.dataset.transform = test_transform

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

# Define loss, optimizer, and scheduler
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
scheduler = StepLR(optimizer, step_size=5, gamma=0.1)

# Track training loss
train_losses = []

# Training loop
for epoch in range(20):
    model.train()
    running_loss = 0.0
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    scheduler.step()
    train_losses.append(running_loss)
    print(f"Epoch {epoch+1}/20 complete, Loss: {running_loss:.4f}")

# Evaluation
model.eval()
all_preds, all_labels = [], []
with torch.no_grad():
    for images, labels in test_loader:
        images = images.to(device)
        outputs = model(images)
        _, preds = torch.max(outputs, 1)
        all_preds.append(preds.cpu().numpy())
        all_labels.append(labels.cpu().numpy())

```

```

all_preds.extend(preds.cpu().numpy())
all_labels.extend(labels.numpy())

# Classification metrics
class_names = full_dataset.classes
print(classification_report(all_labels, all_preds, target_names=class_names))

# Confusion matrix
cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(12, 10))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.title("Confusion Matrix - ResNet50")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Plot training loss curve
plt.figure(figsize=(8, 5))
plt.plot(range(1, len(train_losses) + 1), train_losses, marker='o')
plt.title("Training Loss Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("Total Training Loss")
plt.grid(True)
plt.tight_layout()
plt.show()

```

Epoch 1/20 complete, Loss: 224.2246
 Epoch 2/20 complete, Loss: 103.0149
 Epoch 3/20 complete, Loss: 77.6501
 Epoch 4/20 complete, Loss: 63.3028
 Epoch 5/20 complete, Loss: 47.3332
 Epoch 6/20 complete, Loss: 18.8402
 Epoch 7/20 complete, Loss: 7.1968
 Epoch 8/20 complete, Loss: 5.6177
 Epoch 9/20 complete, Loss: 3.7071
 Epoch 10/20 complete, Loss: 3.5209
 Epoch 11/20 complete, Loss: 2.4319
 Epoch 12/20 complete, Loss: 2.1320
 Epoch 13/20 complete, Loss: 1.7934
 Epoch 14/20 complete, Loss: 1.6472
 Epoch 15/20 complete, Loss: 1.8930
 Epoch 16/20 complete, Loss: 1.9758
 Epoch 17/20 complete, Loss: 1.5551
 Epoch 18/20 complete, Loss: 1.6613
 Epoch 19/20 complete, Loss: 1.8195
 Epoch 20/20 complete, Loss: 1.9083

	precision	recall	f1-score	support
Agriculture	0.98	0.97	0.98	146
Airport	0.93	0.94	0.94	146
Beach	0.97	0.99	0.98	153
City	0.98	0.98	0.98	170
Desert	0.94	0.97	0.95	153
Forest	0.97	0.97	0.97	158
Grassland	0.98	0.97	0.98	162
Highway	0.97	0.96	0.97	161
Lake	0.97	0.98	0.97	178
Mountain	0.96	0.94	0.95	166
Parking	0.99	0.97	0.98	166
Port	0.98	0.97	0.97	174
Railway	0.94	0.93	0.93	149
Residential	0.98	1.00	0.99	169
River	0.94	0.95	0.95	149
accuracy			0.97	2400
macro avg	0.97	0.97	0.97	2400
weighted avg	0.97	0.97	0.97	2400

Confusion Matrix - ResNet50

